

# Generative Adversarial Networks

Sahin Olut

Department of Computer Engineering  
Istanbul Technical University

November 4, 2017

# Outline

- 1 Motivation
  - Why should we study generative models?
  - Some results from recent GAN works
- 2 How does GAN work?
  - GAN Architecture
  - Formulation of GAN
  - Training procedure of GANs
- 3 Applications of GANs
  - Computer Vision
  - Reinforcement Learning

# Outline

- 1 Motivation
  - Why should we study generative models?
  - Some results from recent GAN works
- 2 How does GAN work?
  - GAN Architecture
  - Formulation of GAN
  - Training procedure of GANs
- 3 Applications of GANs
  - Computer Vision
  - Reinforcement Learning

# Motivation

- We can restore the missing data by generative models. (Image inpainting, super-resolution)

# Motivation

- We can restore the missing data by generative models. (Image inpainting, super-resolution)
- We can generate new examples to enhance our classifier networks. (Data augmentation strategy)

# Motivation

- We can restore the missing data by generative models. (Image inpainting, super-resolution)
- We can generate new examples to enhance our classifier networks. (Data augmentation strategy)
- If we want our computers to understand, we have to teach them to create. (I do not understand what I cannot create. – Richard Feynman)

# Outline

## 1 Motivation

- Why should we study generative models?
- Some results from recent GAN works

## 2 How does GAN work?

- GAN Architecture
- Formulation of GAN
- Training procedure of GANs

## 3 Applications of GANs

- Computer Vision
- Reinforcement Learning

# Recent works

- Photo-Realistic Single Image Super-Resolution Using a Generative Adversarial Network [Ledig et al., 2016]

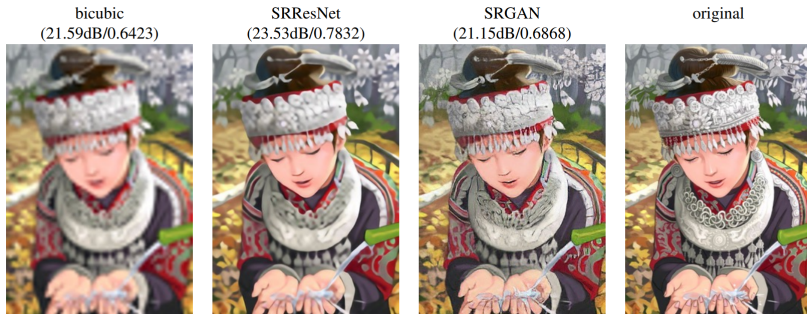


Figure: Work by Ledig et al., 2016



# Recent works

- Plug & Play Generative Networks: Conditional Iterative Generation of Images in Latent Space [Nguyen et al., 2017]

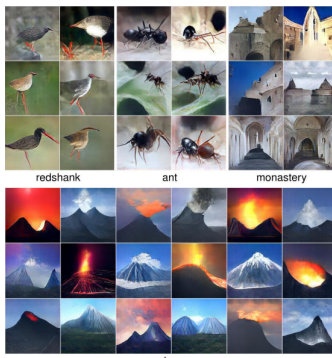


Figure: Synthetic images generated from ImageNet classes.

# Outline

- 1 Motivation
  - Why should we study generative models?
  - Some results from recent GAN works
- 2 How does GAN work?
  - GAN Architecture
  - Formulation of GAN
  - Training procedure of GANs
- 3 Applications of GANs
  - Computer Vision
  - Reinforcement Learning

# Discriminator and Generator Networks

- What is a generative model?

# Discriminator and Generator Networks

- What is a generative model?
- Discriminator's role in GAN is to predict whether the input is generated or sampled from training data.
- The aim of generator is to capture the distribution of the training data.
- According to Goodfellow et al., it is a minimax game between generator and discriminator where generator tries to fool discriminator.(There are some debates about it.)

# Outline

- 1 Motivation
  - Why should we study generative models?
  - Some results from recent GAN works
- 2 How does GAN work?
  - GAN Architecture
  - **Formulation of GAN**
  - Training procedure of GANs
- 3 Applications of GANs
  - Computer Vision
  - Reinforcement Learning

- The description of GANs leads us to formulation for loss:

$$\min_G \max_D V(D, G) = \mathbb{E}_{x \sim p_{\text{data}}(x)} [\log D(x)] + \mathbb{E}_{z \sim p_z(z)} [\log(1 - D(G(z)))] \quad (1)$$

- Where both networks rely on gradients flowing through discriminator.

# Outline

- 1 Motivation
  - Why should we study generative models?
  - Some results from recent GAN works
- 2 How does GAN work?
  - GAN Architecture
  - Formulation of GAN
  - Training procedure of GANs
- 3 Applications of GANs
  - Computer Vision
  - Reinforcement Learning

---

**Algorithm 1** Minibatch stochastic gradient descent training of generative adversarial nets. The number of steps to apply to the discriminator,  $k$ , is a hyperparameter. We used  $k = 1$ , the least expensive option, in our experiments.

---

**for** number of training iterations **do**

**for**  $k$  steps **do**

- Sample minibatch of  $m$  noise samples  $\{z^{(1)}, \dots, z^{(m)}\}$  from noise prior  $p_g(z)$ .
- Sample minibatch of  $m$  examples  $\{x^{(1)}, \dots, x^{(m)}\}$  from data generating distribution  $p_{\text{data}}(x)$ .
- Update the discriminator by ascending its stochastic gradient:

$$\nabla_{\theta_d} \frac{1}{m} \sum_{i=1}^m \left[ \log D(x^{(i)}) + \log \left( 1 - D(G(z^{(i)})) \right) \right].$$

**end for**

- Sample minibatch of  $m$  noise samples  $\{z^{(1)}, \dots, z^{(m)}\}$  from noise prior  $p_g(z)$ .
- Update the generator by descending its stochastic gradient:

$$\nabla_{\theta_g} \frac{1}{m} \sum_{i=1}^m \log \left( 1 - D(G(z^{(i)})) \right).$$

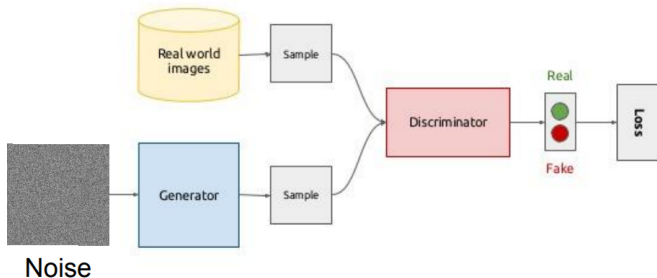
**end for**

The gradient-based updates can use any standard gradient-based learning rule. We used momentum in our experiments.

---



# GAN Framework



From now and on, we have a basic grasp of GAN, therefore we can code our own GAN! The starter code can be found in my GitHub Repository: [github.com/norveclibalikci/InzvaGanStarter](https://github.com/norveclibalikci/InzvaGanStarter)

# Outline

- 1 Motivation
  - Why should we study generative models?
  - Some results from recent GAN works
- 2 How does GAN work?
  - GAN Architecture
  - Formulation of GAN
  - Training procedure of GANs
- 3 Applications of GANs
  - Computer Vision
  - Reinforcement Learning

# Computer Vision Applications

- Image generation(Plug and Play GAN)
- Style transfer(CycleGAN)
- Image inpainting, super-resolution(SRGAN)
- Image to text (Image captioning - Generative Adversarial Text to Image Synthesis)

# Computer Vision Applications

- Image generation(Plug and Play GAN)
- Style transfer(CycleGAN)
- Image inpainting, super-resolution(SRGAN)
- Image to text (Image captioning - Generative Adversarial Text to Image Synthesis)
- There are many applications which are not covered above.

# Outline

- 1 Motivation
  - Why should we study generative models?
  - Some results from recent GAN works
- 2 How does GAN work?
  - GAN Architecture
  - Formulation of GAN
  - Training procedure of GANs
- 3 Applications of GANs
  - Computer Vision
  - Reinforcement Learning

# Applications to real life

- Discriminator can be rewarded for labeling correctly and a new loss can be defined by that way. (Still on research)
- Generating environment and test for reinforcement learning applications.
- Simulating particle experiments like they do in CERN.

Many things can be done with GANs however, GANs have limitations as well.

- There is no training procedure that has been proven to be successful.



- There is no training procedure that has been proven to be successful.
- Sometimes discriminator learns faster than generator(predicts everything correctly), which leads a gradient problem to generator. In some cases, it is vice-versa as well.

- There is no training procedure that has been proven to be successful.
- Sometimes discriminator learns faster than generator(predicts everything correctly), which leads a gradient problem to generator. In some cases, it is vice-versa as well.
- After from some point, generator keeps generating similar examples.

- There is no training procedure that has been proven to be successful.
- Sometimes discriminator learns faster than generator(predicts everything correctly), which leads a gradient problem to generator. In some cases, it is vice-versa as well.
- After from some point, generator keeps generating similar examples.
- Losses of models are not meaningful as classifier's. It just keep oscillating back and forth.

# Loss graphic of a GAN



Figure: Taken from <http://www.rricard.me>

# Summary

- GANs are **very** powerful tools for generating new samples from data yet it has serious issues.
- GAN uses **semi-supervised** approach therefore, there is no need for data-labeling.
- **Unsupervised learning is the cake of true AI.** (Yann LeCun, Head of Facebook Research)

# For Further Reading I



Chintala et al.

GAN Hacks

[github.com/soumith/ganhacks](https://github.com/soumith/ganhacks)



Goodfellow et al.

Generative Adversarial Networks

*arXiv*: 1406.2661, 2014.